

GPT-6 Astra: an automated AI Engineer you can hire for <\$6 an hour

We spent 20B+ tokens of GPT-6 Astra to explore everything. Here's our learnings.

SEP 04, 2026

GPT-6 Astra, the first [Stargate](#) and [lightly looped](#) supermodel from OpenAI, [launched today](#), cleanly beating [Table 5.1](#) on many metrics including completely saturating the hardest versions of [FrontierMath](#) (97.6%) and [ARC-AGI-3](#) (99.9%). Lots of demos will focus on typical talk tracks like the [computer use](#) to the [Pokemon playing](#) to [Blender](#) to the [scientific](#) and [cybersafety](#) benchmarks ([system card](#)). Greg says [AGI is here](#), and Jakub says it is [finally the Automated AI Research Intern](#) he wanted.

We aren't qualified to talk about those, but we got early access and threw it at every practical, real-life task we could think of. After burning **over 20B tokens of Astra**, we can confirm the most surprising finding: **GPT-6 Astra is one of a new class of models¹ that are fully capable AI Engineers in their own right**. They now help you **choose and train models, label data** (both helping you label and then using your labels for active learning, like [SAM](#)), **keep pipelines saturated, instrument and read logs, deploy and debug entire systems** in one shot, fan out and **command and eval subagents** (including agents running other models), and keep coherence over **billions** of tokens of a single agent thread.



Raising Your Ambitions

We've written before about [the high-return activity of raising your aspirations for LLMs](#). **Our experience has made us exponentially more ambitious than we have ever been.** Over the past month, we went from prompting humans for a fun “[Kill My SaaS](#)” competition², to building a [dozen internal/personal tools](#), including [4 previously paid SaaS tools](#), fully [redesigned my personal site](#), made an incomplete but functional [replacement of GitHub + Vercel](#), trained game AI for [a strategy board game with 10,000x more legal moves than Go](#), saved tens of thousands of dollars in personal finance cleanups, [republished my old book](#) with synced audiobook audio and printed physical editions, and even [more ambitious](#) projects we will launch soon.

The \$6 an hour number might sound surprising, but that's exactly what we saw in [our testing](#) - 33 tokens per second at a max \$50 per million token rate. Given that Astra is more token efficient than Sol and Fable ([independently confirmed by Artificial Analysis](#)), it often means that Astra is simultaneously also the best fast-and-smart model you can buy (assuming our preview latency holds for GA), outside of [Spark 1.3](#).

Run telemetry



Compare successful-call latency, token use, and reliability. Failures and measurement coverage remain visible.

[Collection](#)
[Extraction](#)
[Repair](#)
 Model filter applies to collection. Each stage is a different workload.

Aborted Gemini 3.8 Flash, ZCode and Deep Code runs are excluded from these comparisons. [Run exclusions](#) ↗

Latency

Median wall time per successful call. Compare tail latency in the p95 column below.

Muse Spark 1.3		75 s
	132/132 successful calls measured · muse · high · native	
GPT-6 Astra		99 s
	132/132 successful calls measured · codex · low · live	
Grok 4.6		158 s
	133/133 successful calls measured · cursor · - · native	
GPT-5.6 Sol		158 s
	132/132 successful calls measured · codex · low · live	
Claude Opus 5		213 s
	133/133 successful calls measured · claude · high · native	
SWE-1.7		238 s
	131/131 successful calls measured · devin · - · native	
Claude Fable 5.1		311 s
	131/131 successful calls measured · claude · high · native	

Token use


Median reported tokens per successful call. Input includes cached tokens; shorter output is not a quality score.

Output tokens	Input tokens		
GPT-6 Astra		1,192	
	132/132 successful calls measured · codex · low · live		
GPT-5.6 Sol		2,605	
	132/132 successful calls measured · codex · low · live		
Muse Spark 1.3		4,170	
	132/132 successful calls measured · muse · high · native		
Claude Fable 5.1		5,368	
	131/131 successful calls measured · claude · high · native		
Grok 4.6		6,900	
	133/133 successful calls measured · cursor · - · native		
Claude Opus 5		8,914	
	133/133 successful calls measured · claude · high · native		
SWE-1.7		14,871	
	131/131 successful calls measured · devin · - · native		

[see logs](#)

Managing fleets of subagents (individually tweaked, bounded concurrency)



Now of course, if you just throw on Astra at Ultra you're gonna burn through a lot more than \$6 per hour.... because it is so dang good at parallelizing. Depending on the task in practice we were often ramping up **between 20-50 agents in parallel**, of course all managed by one main Astra agent.



Monitoring its own runs, starting and stopping waves

This is basically what you would pay a junior AI Engineer to do — babysitting runs, staring at data, finding issues, fixing, rerunning, ad infinitum. You could hire someone at \$200-\$1000 a day, or you can hire GPT-6 for \$100 over 2 days to do this.



GPT-6 Astra: an automated AI Engineer




Making model benchmarks, handling budgets, making estimates, scaling up runs, getting human ratings

Because of course you need all these capabilities to run your own AI engineering program, because of course OpenAI already uses GPT-6 to do this internally...

GPT-6 Astra: an automated AI Engineer



GPT-6 Astra: making custom evals and writing report



example [here](#)

Or you can get Astra to trivially whip up your own **personal Arena.ai clone** for tuning your prompts, picking models for your task, or aligning your own preference model!

WORKBENCH

Prompt builder

Source benchmarks

Image → prompt

Rate results

Generation runs

Saved tests

AI prompt review

Model pricing

Settings

Workspace / Rate results

Refresh

YOUR WORKSPACE · HUMAN PREFERENCES

Rate results

Export generations & votes

Pick the better result. Every vote stays local.

Pairwise arena

Global ratings & refusals

Automated judge

aligning pref model

Global human ratings

Across included workspace generations; is

0 generations and 0 votes excluded. Resto

Bradley-Terry · all votes fitted together, inde


excluded; unvoted models are unranked.

How scores are calculated · 28 comparisons across 1 prompts

Model / profile	BT score	Corrects - W/L/T	Both bad	Successful	Refusals / submitted	Errors / pending
Reve 2.1 f13951f · pool 182702	1605	7 · 7/0/0	0	4/4	0/4 · 0%	0/0
Grok Imagine a6c35c46 · pool 182702	1379	7 · 6/1/0	0	3/3	0/3 · 0%	0/0
Grok Imagine · Quality 18270261 · pool 182702	1196	7 · 5/2/0	0	1/1	0/1 · 0%	0/0
Nano Banana 2 1bdf133 · pool 182702	898	7 · 3/4/0	0	6/6	0/6 · 0%	0/0
Qwen Image 3.0 Pro 8e8dabc9 · pool 182702	898	7 · 3/4/0	0	4/4	0/4 · 0%	0/0
Nano Banana 2 Lite 23b8582a · pool 182702	767	7 · 2/5/0	0	3/3	0/3 · 0%	0/0
Seedream 4.0 fa080a6c · pool 182702	629	7 · 1/6/0	0	3/3	0/3 · 0%	0/0
Seedream 5.0 Pro c8628a06 · pool 182702	629	7 · 1/6/0	0	3/3	0/3 · 0%	0/0
FLUX.2 Dev 30082f8c	Unranked	0 · 0/0/0	0	1/1	0/1 · 0%	0/0

vibecoded personal arena to tune models

The overall conclusion you should have is that **OpenAI have clearly trained a model that is capable of automating much of their own AI Engineering**, and it is finally time that you learn to exploit Astra- and Fable-class models and be far, **far more unreasonable** with your own expectations of what you can do with agents now.



-
- 1 We are [running similar work](#) on Grok, Fable and other similar frontier models but OpenAI was most generous with trial limits so this gets the writeup - but the agentic coding patterns discussed here will likely apply to all such late 2026 frontier models.
 - 2 Many of you are waiting to hear results... sorry for the radio silence! we got... busy! We will announce winners and reimbursements and best attempts.